DroiDTN: Optimizing Bandwidth over Wifi and 3G on Smartphones

Shreshth Singhal

May 6, 2013

Abstract

Lorem ipsum dolor sit amet, consectetur adipisicing elit, sed do eiusmod tempor incididunt ut labore et dolore magna aliqua. Ut enim ad minim veniam, quis nostrud exercitation ullamco laboris nisi ut aliquip ex ea commodo consequat. Duis aute irure dolor in reprehenderit in voluptate velit esse cillum dolore eu fugiat nulla pariatur. Excepteur sint occaecat cupidatat non proident, sunt in culpa qui officia deserunt mollit anim id est laborum.

1 Introduction

Worldwide, the use of smartphones is becoming ubiquitous. One of the primary advantages of using smartphones over their "non-smart" counterparts is the ability to access the internet, either through wireless (wifi) or cellular data networks (3G, 4G LTE etc.). More and more smartphones are available that are leveraging these networks to bring more and more content to users - be it through apps, services, or directly browsing the internet.

1.1 Motivation: comparison of wifi and 3G

As smartphones have proliferated, so have cellular data networks. We know these through various names - 3G, 4G LTE etc. (and we will refer to these collectively as 3G for short from now on). Ericsson estimates that there are 6.3 billion mobile subscriptions among 4.4 billion people worldwide, with international wireless penetration reaching 89%.[6] To understand this staggering growth, it is important to realize that only 7 years ago, there were only about 2.5 billion subscriptions, and an average penetration of 40%.[7] However, this immense growth has imposed infrastructural limits on data networks because of which users do not have access to unlimited network usage. We see this in the way that data plans offered by network providers are moving away from the concept of "unlimited" plans towards expensive plans, data caps and throttling.[4][14] Further, data networks are not available uniformly everywhere. There are plenty of pockets where data connectivity is poor or nonexistent.[3][13] Data networks also have an associated cost as reflected in data plans. This is usually a package cost up to a certain limit of data usage, and a per-byte fee thereafter. Data networks have also been shown to be less power-efficient than wifi when transferring data, due to lower transfer rates (although the power usage rate over time is around the same).[9, p. 8][1, p. 249][5, p. 280]

1

On the other hand, wifi tends to be free, or at the very least, cheaper than data networks. wifi networks also usually don't have aggressive caps or thresholds on their usage, and don't face data throttling problems. However, it is important to note that wifi is not ubiquitous either. Wifi routers are stationary and limited in their range, considerably more so than data networks which are provided through long-range cellular towers. Because of this, wifi actually uses more energy when idling, as compared to more powerful cellular networks.[5, p. 280] [11, p. 1] Further, installing wifi infrastructure is still quite expensive.[12, p. 105] Many people face the problem of not being able to connect to wifi on the move. An even more annoying problem is when our phones connect to a wifi network but are unable to send or receive data at a usable rate because of poor connection.

Therefore, we can see that neither wifi nor 3G can function perfectly on their own. They have differing costs, availability and bandwidth, which themselves differ across different locations and times. The ideal goal would be to provide the user with good, usable network connections as often as possible, in a manner that reduces the cost to them and provides them with a good experience. Therefore, smartphones should be capable of switching between wifi and data networks smartly in a manner that minimizes cost, while at the same time ensuring the quality of service is good.

1.2   Problem - Network Scheduling

We have seen that wifi and 3G do not work best alone. Our goal therefore is to leverage them in combination in a way that makes the user happy. The characterizations of wifi and 3G above are an indication of the direction we should take to be able to provide a better experience for the user - we need to focus on **quality of service** and **limiting cost**.

As we have seen before, 3G tends to be more expensive than wifi. Offloading network traffic to wifi whenever possible is hence an ideal goal. It may seem trivial to do so - just use wifi when it is available and 3G otherwise. However, there are improvements we can make to this trivial solution. We can leverage the fact that certain network traffic on smartphones is **delay-tolerant**. This means that we can delay 3G traffic in the hopes of waiting for wifi to show up, as long as it doesn't affect the user's experience. An example of this is the case of e-mail syncing. E-mail clients on phones constantly synchronize webmail in the background. Oftentimes, however, this process can be delayed, since the user does not check email all the time, and would not notice the difference if email came in a minute later.

There are complications that arise because we want to provide the user good quality of service. Because of this, we obviously cannot delay traffic immediately affects the user - for example, if the user is watching a video or clicks on a link on a website, we cannot delay it since the user is directly interacting with the phone and cares about getting results immediately. Further, we also do not want to offload to wifi in cases where the wifi signal is poor and would result in a bad experience for the user.

Note that we are focusing here on traffic that is produced by the phone, i.e. network traffic leaving the phone. This is because, in most common applications, even if the phone needs to receive some information, it is not done unilaterally by the server. Rather, the phone makes a request to the server for information, and the server responds to this request. For example, email syncing is done by the phone pinging the webmail server to find out if any new email has been received. This simplifies the model since we only need to focus on the network traffic generated by the phone, and can abstract away the traffic received by the phone as responses to the phone's requests, and hence part of the same network transaction.

Therefore, our **problem statement** resolves to the following: we need to schedule network traffic generated by the phone through either wifi or 3G in a manner that minimizes cost, while at the same time ensuring good quality of service. In effect, this is a **network scheduling problem**. To this end, we can do many smart things, such as leveraging delay tolerance of some network traffic to avoid using 3G and hence reducing cost without affecting the user.

2 Overview of solution

2.1 Leveraging delay tolerance - theoretically

The most nebulous part of the solution as described above is the notion of **delay tolerance**. Network traffic is delay tolerant if it can withstand some delay to its transmission without affecting the user. Note that not all delay-tolerant traffic is equal - some applications' traffic can be delayed more than others. For example, email syncing can be delayed for times on the order of minutes, since users tend to check for new email infrequently. [CITE?] On the other hand, video streaming applications have a much shorter threshold - we can only delay traffic as long as the video frame-buffer is not close to being empty, since the user would be affected if the buffer became empty and the video stopped playing. Delay tolerance is therefore defined per-application and is a measure, in seconds, of how much that application's traffic can be delayed.

This gives rise to a few questions. How exactly do we tell what network traffic is tolerant to delay? How much should we delay such traffic in the hopes of getting wifi soon? How do we quantify the "hope" of getting wifi soon? My method relies on the theoretical framework developed by Yetim et al., as described in Section 3, which requires perfect knowledge about two things:[15]

1. *network traffic*: How much network traffic is going to be generated by the phone for all times in the future? What is the nature of this network traffic - how much of it is tolerant to delay, and how much must be sent immediately? What is the individual delay tolerance of each of the applications that are producing network traffic?

2. *availability of wifi and 3G*: What is the nature of wifi and 3G that the user (and the phone) is going to experience at all times/all locations? What is the speed, signal strength and availability of these networks?

It is easy to see that with perfect knowledge of these things, we can schedule wifi and 3G in an optimal manner. We know exactly the sort of traffic we are going to see, and we know the exact nature of the networks that we are going to be sending them over. We can therefore schedule traffic that is not delay tolerant over whatever network is available at the time the traffic needs to be sent. For traffic that can be delayed, we can gauge whether it is possible to wait and send it over wifi, and still remain within the limits of its delay tolerance. If we cannot do so, we can just send it over 3G.

2.2 A real solution - machine learning

In a real system, we are not going to be presented with this perfect information. Therefore, we must learn this from past measurements, in essence reducing this to a machine learning problem. Corresponding to the two aspects of perfect information required (from the previous section), we have two sources of past information we can learn from, so as to emulate this perfect information:

1. *application delay tolerance estimation*: We can track the past network usage of each application on the smartphone. From this, we can learn the delay tolerance of the app, for example, by looking at the time the app usually takes between sending successive requests. This will give us an approximate handle on how long the app's traffic can be delayed.

   Then, when we are actually scheduling traffic, we can look at which app generated the traffic, and use its delay tolerance estimate to figure out if we can delay it or not.

2. *future wifi prediction*: We can track the wifi and 3G availability and quality at different physical locations as the user of the phone moves around. Then, when we are scheduling traffic, we can look at the user's current location and direction of motion to predict how long it will take them to be in range of usable wifi. If the traffic is delay tolerant enough to wait for this time, then we can delay the traffic, otherwise we can send it immediately over 3G. Of course, if our prediction of the time required to reach wifi was incorrect, we need to failover and immediately send the traffic over 3G.

   Note that this method will be more successful the more dense and accurate our mapping of wifi and 3G in different locations is. Therefore, it would make sense to share this data across all users so as to collaboratively have a better understanding of the network map.

2.3   Putting it all together

After gaining a high-level understanding of our problem and possible insights on tackling it in the previous sections, we can see that there are multiple pieces that must be built:

- **Policy**: For every network packet that we see, how much should we delay it? This is what we will refer to as the "policy" for delaying traffic. This policy is determined by the two different measurements that we described in the previous section:

  1. An *application delay tolerance estimator*: This estimator tracks the network usage statistics of each application on the smartphone, to learn the delay tolerance of each app.
  2. A *time-to-wifi predictor*: This predictor tracks wifi and 3G signal strength and availability at different places as the user walks around with their phone. From these, it can learn to predict the time it will take to get in range of wifi, given the user's current location and direction of motion.

  Using these measurements, we can determine a policy to delay traffic. One that immediately springs to mind would be to delay traffic if

  $$\text{delay tolerance} \geq \text{time-to-wifi}$$

  However, it is important to note that this is not the only sensible policy. Other policies like a fixed delay or zero delay for all traffic may also make sense in certain cases. For example, if our map of wifi and 3G in different locations is very sparse, then it is likely that our time-to-wifi predictor will be inaccurate. Because of this, it may make sense to replace the time-to-wifi value with infinity or double each value just to be safe. Equivalently, if we don't have much confidence in our application delay tolerance estimator, then we can just replace the delay tolerance values with zero or some small fixed value just to be safe.

- **Mechanism**: Now that we have a policy for when to delay network traffic, and how long to delay it, we still need to set up the mechanism to make the delaying work. This "mechanism" refers to some actual piece of code that will run on the smartphone and delay packets of data for some time, depending on the policy. This mechanism could be run in many ways, in the application level, or within the OS. We want to try and ensure that the mechanism runs in a way that makes it easy to program and invisible to the user. We also want to ensure that we abstract away the complexities as much as we can from application developers who are going to be writing apps which generate network traffic.

Therefore, the pieces of the solution as described are:

Figure 1: Method

BELOW IS WORK IN PROGRESS

2.4   Setting

I will be working with an Android smartphone - specifically a Galaxy Nexus running Android 4.1.1 (Jelly Bean API 16). The setting for my experiments is the Princeton University campus. The campus is characterized by wifi that is open and available freely in all buildings and also in a lot of open spaces. However, there are certain stretches which do not have wifi or suffer from poor connectivity. 3G is available across campus on the ATT network which I will be using for testing.

This small geographic area therefore has a wide array of wifi profiles, and the prevalence of 3G allows it to be catchall or a fallback for network traffic in places where wifi is not available. This makes it a good location for testing our system - we can easily test if the app aggressively delays traffic in the places where wifi is not available, since it is quite likely that the user will be in range of wifi soon (because wifi is quite widely available).

I will be testing with an open source email client for Android, called K-9. This is the app that I will be testing with for determining delay tolerance, and testing the policy. It will also be used for data collection

2.5   Approach

There are multiple approaches which can be taken to implement the policy and mechanism as described above.

- *Application-level*: One way would be to build a system entirely on the application-level, without going into the kernel/OS. Firstly, we would be to build the *time-to-wifi predictor* as an app for the smartphone, which would store all the information about wifi and 3G at different locations, and learn from this data to predict time-to-wifi.

  Each other app, before making a network request, would ask this app for how much time it will take to be in range of wifi. Then the apps can internally make judgements, based on their own known delay tolerance, about whether to delay their traffic or not. We therefore place the *delay tolerance estimator* in the apps themselves, because they are likely to have the best idea about their own delay tolerance. The *policy* and *mechanism* are also implemented in the apps themselves.

  Therefore, the time-to-wifi predictor is in a new central app, and the existing apps act as their own delay tolerance estimators and their own mechanisms for delaying traffic.

  The major advantage of this method is that it can be implemented entirely in the application layer without requiring the OS to support it. Therefore, we can write the app ourselves, and we can then design apps to interact with it. The entirety of the solution is therefore encapsulated among applications without going to the OS, making it easy to write and distribute. Another great advantage is the fact that this solution ensures that the the app developers are aware of the delay tolerance heuristic. Because of this, they can provide the delay tolerance value of their app, rather than needing us to guess it. Further, the mechanism is being implemented in the apps themselves, and this ensures that delays are graceful. Instead, if the app was not expecting to see such delays, it may respond in strange ways, like spinning off multiple threads to send the same network traffic.

  However, it is obvious to see that it requires modifications to all other applications. Application developers need to add code to talk to the time-to-wifi predictor, need to calculate

their app's delay tolerance, and also need to manage the delaying of network traffic. This makes this solution less appealing - the development process is harder, and it requires collaboration with application developers.

- *OS-level*: Alternately, an approach could be to implement the system in the OS itself, thereby taking the onus out of the hands of the developers. The OS has easy access to all the statistics required by the policy maker, namely apps' network usage and wifi and 3G locations. We therefore, would only be required to store this data and implement machine learning algorithms to leverage them. This would make the OS capable of being the *time-to-wifi predictor* and the app *delay tolerance estimator*, i.e. the *policy* is implemented entirely in the OS.

  Then, whenever the OS gets a network request, it can check which app made that request. Now, the OS can use the policy to find out how much to delay that network request. If we make changes in the network stack, the OS can also delay the network request. We could implement a data structure to store all the network packets that are being delayed, sorted by timestamp of when they need to be necessarily sent, and use this to delay traffic till its time expires, or until we reach wifi. Therefore, the *mechanism* can be implemented in the OS also.

  This abstracts away the complexity from application developers, and allows a clean encapsulation of the entire system in the OS. Further, since all system operations go through the OS, we get much better accuracy in storing statistics. For example, storing network statistics is much more accurate if we store them every time the OS gets a request for some network activity, rather than if we poll the combined network statistics every so often.

  On the other hand, writing code in the OS is cumbersome to say the least. Current versions of the Android system have over a million lines of code, and it is hard to parse and figure out where to insert our code to modify the behaviour. Further, most applications are unaware of the delays that are being introduced. They may react in strange manners when they see that their requests are not receiving any response - something like repeatedly send their requests, or maybe even crash.

Our approach is an amalgam of the above two methods. We tried to limit the burden on application developers as much as possible. At the same time, for ease of programming and distribution, we tried to avoid descending to the OS-level. To give a broad overview:

- The *time-to-wifi predictor* will be entirely outside the OS - on the application level, and offline on remote servers. This will require using Android APIs to collect the statistics about wifi and 3G at different locations. We expect this to reduce the accuracy of data collection, compared to collecting the same data in the OS. Our method will poll for these statistics every so often, whereas the OS will have a much more continuous view of how wifi and 3G change over time. However, we believe that this error will be negligible compared to the complexity of implementing the same system in the kernel. Note that in either case, the application developer and the end-user remain uninvolved, which is what we want.

  Using the data collected, we will perform machine learning offline, on a remote server, and use this to predict time-to-wifi.

- The *delay tolerance estimator* will also be outside the OS. However, it will not be in the applications themselves, as we had described above. While such a system would have

allowed the applications to give exact values of their delay tolerances, it would also involve the application developer, which is something we want to avoid.

Instead, similar to the time-to-wifi predictor, we will develop an app that collects data about other apps' network usage, using Android APIs. It faces the same problem of coarse granularity as the time-to-wifi predictor, since we are polling for these statistics rather than logging network traffic as it happens (which would be possible in the kernel). However, since we can poll at a high frequency here, we actually do not expect the error to be that much.

After collecting this data, we can perform machine learning on it. In this case, we can do the learning on the phone itself, and use this to estimate an application's delay tolerance.

- The *mechanism* for delaying network traffic needs to be implemented in the OS. Implementing this in the application layer implies the involvement of application developers to delay the network traffic of their own application. We cannot build an app that will delay the network traffic of other apps - this would be a security issue.

  Hence, the only other way to implement the mechanism is to intercept network traffic in the OS, in the network stack. We can then predict the delay tolerance of the traffic, and delay it if necessary, depending on the time-to-wifi estimate (these can be found by talking to the apps described above). All of this can be done invisibly to the application developer, but requires tinkering with the network stack.

3 Related work

3.1 Offloading data onto wifi

The problem of offloading 3G data onto wifi is not a new one. Much has been written about it, and many attempts have been made to figure out better ways to do it. A paper by Yetim and Martonosi in CHANTS '12 talks about a theoretical framework for optimizing network traffic on smartphones between wifi and 3G. The framework requires perfect knowledge about all network data streams on the smartphone, and about the availability of wifi and 3G in the future. Given this knowledge, they resolve the problem to a mixed-integer linear programming problem, where they try to minimize the cost of traffic sent over 3G, subject to the constraints defined by the data streams and network availability. The linear program leverages delay tolerance and assigns the optimal network to be used by each data unit in order to minimize this cost.[?] This method makes intuitive sense - they have perfect knowledge about how much data the apps are going to generate in the future, and also about the wifi and 3G characteristics at all those times in the future. They also assume knowledge of delay tolerance of the network data. Given all this knowledge, they can schedule all network data over wifi and 3G in a manner that minimizes cost. Their method is optimal in a theoretical sense, and allows insight into the possible ways of implementing a real solution. However, it is not a realistic solution in itself, since solving the required linear program is too slow on the phone. Further, they assume a lot of knowledge about the real world. Our goal will be to try to implement an efficient, standalone system, based on their theoretical framework. This means that we will be implementing a network scheduling system that leverages delay tolerance, and uses machine learning techniques to emulate the perfect knowledge of the world.

Balasubramanian et al. also talk about offloading of data onto wifi. Their goal is to try to reduce energy cost of sending data over 3G. They similarly develop an online algorithm that leverages delay tolerance of applications to schedule network traffic in a manner that reduces the high energy cost of sending data over 3G.[5]. Their method is quite different from ours since their goal is to minimize energy usage, rather than cost over 3G. Because of this, they may prefer to send some network data over 3G rather than wifi if it will allow them to exploit the high energy cost of ramping up the 3G signal. This is very different from our approach which minimizes cost and will always prefer usable wifi over 3G.

Another paper by Lee et al. talks about the real-world performance of wifi offloading. They conclude that wifi offloading without taking into account delay tolerance is quite effective in itself, but can be made more efficacious by adding delay (they add a constant delay to network traffic). They further propose a framework similar to the above to generate network traffic and measure the effect of offloading under delay.[9] Ra et al. form a similar theoretical framework to take advantage of delay tolerance. However, in their case, the objective function minimizes energy usage rather than 3G cost. They delve into the differing energy requirements of wifi and 3G, and also across various wifi networks, and try to delay traffic until a more energy-efficient network is available.[10]

3.2   Predicting time-to-wifi

Ananthanarayan et al. discuss ways of predicting time-to-wifi. Their goal is to reduce the number of unnecessary wifi scans done by the phone, to preserve energy. They delay scanning for wifi by predicting when wifi is expected, so as to reduce energy usage spent on idle scans. They use bluetooth patterns and cellular tower information for making the prediction. This is based on the assumption that users regularly come into contact with a few often-used bluetooth devices (like a mouse or a mobile device/headset), and a few cellphone towers.[1] As we see below in Section 4, these variables are very different from the ones that we use for predicting time-to-wifi. Similar work on predicting wifi and 3G availability was also done by Han et al. [8]

Rahmati et al. have the same goal - preserving energy. Wifi has greater power consumption when on and idling, but is more efficient than 3G when transferring data. Therefore, they try to keep the wifi interface off as long as possible, until they are in a region where wifi is highly probably to be present. Then, they use wifi to make transfers that had previously been delayed, hence saving energy.[11] Once again, since the goal is conserving energy rather than reducing cost, this system behaves very differently from the one we developed. For example, their solution will only turn wifi on when it is highly likely that it is present, whereas our solution tries to opportunistically use wifi as often as possible.

3.3   Predicting delay tolerance

3.4   Mechanism

Todo:

- Expand this with new papers

- Also include overview of Serval and more in-depth look at Ozlem's since these relate to my work greatly

4 Time-to-wifi predictor

The first piece of the puzzle that we consider is the time-to-wifi predictor. To recapitulate, we have a situation where a user is walking around, and may or may not be in range of wifi at different times. This predictor will store the wifi and 3G network characteristics at different times and different places as the user walks around with the phone. Then, given the user's current location and direction of motion, it uses machine learning to predict how much time it will take for the user to be in range of wifi.

4.1 Data collection

To begin with, we need to have a way of collecting data about wifi and 3G at different locations. For this, I built an app that runs on the phone in the background and records wifi and 3G characteristics periodically as the user is walking around, along with the associated locations.

The app uses the `LocationManager` to get a `LocationListener` (these classes and managers are provided in the Android SDK).[2] This is used to listen for location changes. Every time the location changes, we get a new `Location` object, which holds the latitude and longitude of the location, as well as the direction and speed of motion of the user at the time. We then make a recording of the wifi and 3G characteristics at that location as follows:

- The app uses the API provided by `WifiManager` to perform a wifi scan and get `ScanResults`. These `ScanResults` encapsulate information about the wifi characteristics at the location and time where the scan was performed - SSID, signal strength etc.

- It uses another API provided by `TelephonyManager` to find the cellular signal strength, i.e the signal strength of 3G.

Now, since we need to make time-to-wifi predictions, we need to store some notion of how much time it takes to get to wifi from a given location by moving in a certain direction. Therefore, for each record, we want to store not only the wifi profile at that time, but also the wifi profiles in the future. In our case, we store values for 10 time steps; this means that we store, for each timestamp, the wifi experienced at that time step, and the wifi experienced the subsequent 9 time steps. We consider one time step to be one minute, meaning that we store 10 minutes worth of wifi characteristics with each record.

This is basically achieved by putting each each new data-point into a temporary queue, and adding new values to its wifi characteristics for 9 subsequent time steps. Thereafter, it can be removed from the queue and stored in the phone's local store.

Ultimately, then, the app stores a timestamped record containing 3G characteristics at that time, and the wifi characteristics at that time and at 9 future time steps - a triplet containing:

1. location (latitude and longitude), bearing $(0-360°)$, speed of motion (in meters per second)

2. wifi characteristics (signal strength) for next 10 time steps
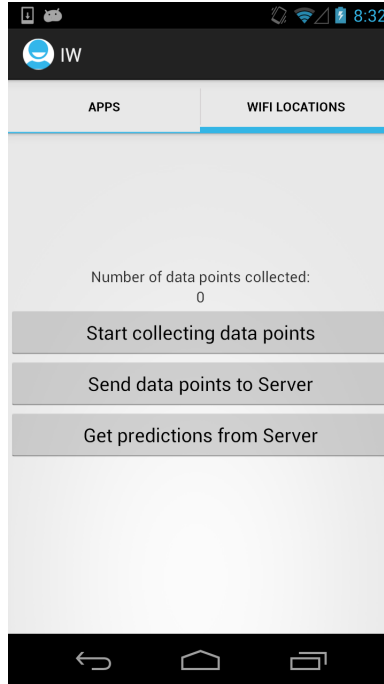
3. 3G characteristics (signal strength)

-

Figure 2: App for collecting wifi and 3G data

4.2 Server

At this point, the phone local store has stored values for the wifi and 3G network characteristics at different locations. Now, there is a question of where do we implement the machine learning algorithms that will leverage the data and predict the time-to-wifi. One solution would be to implement the algorithms on the phone itself. However, the phone's computation capacity may not be sufficient for machine learning algorithms - at any rate, it will be slower on the phone than a computer. It therefore makes sense to do it offline on a centralized server. This has the happy side-effect that it allows us to collaboratively use all users' collected data and hence get a better, denser mapping of wifi and 3G.

Therefore, periodically, the values from the phone's local store are uploaded to a server. This is a server that I wrote on Google AppEngine, which basically has a script to receive data-points from the phone and store them in the very convenient AppEngine database. The app on the phone, when it is in range of wifi, pulls out data-points that it had stored in its local store, and makes a get request to this script to send data points to the server database.

4.3 Preprocessing of data

Some steps were taken to ensure that the data was appropriate for analysis. We filtered out all data points that were not accurate to at least 25m, since we expect wifi availability patterns to change significantly over that distance. Similarly, we filtered out data points that had speed of movement close to 0 m/s. The rationale for doing so is that if a user stood still at the same point for a while and then obtained wifi after some time, that time is not representative of how far that point would be from wifi if the user was walking. Lastly, we experimented with different locations and obtained a cut-off point for the wifi signal strength at which the internet connection is "good
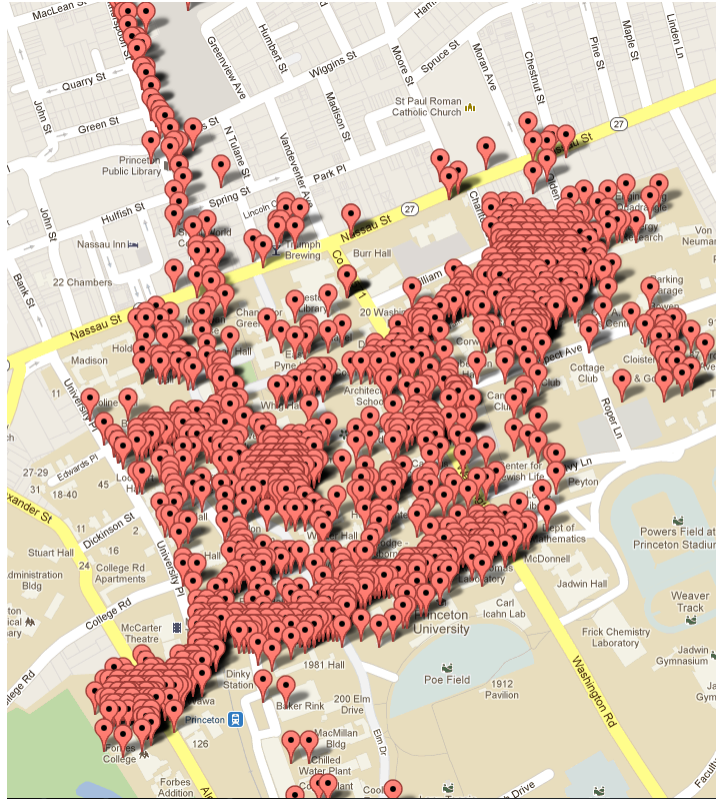
Figure 3: Collected data on the server - `http://cos598b.appspot.com/map`

enough". This solves the quality of service issue, and ensures that the wifi is of usable quality. Based on this cut-off point, and the wifi signal strengths for each data point, we recorded how long it took for each data point to obtain a strong enough wifi signal. We marked the time-to-wifi for points that did not obtain wifi within the 10 minutes as 10 minutes, for ease of analysis.

4.4 Learning algorithms

The AppEngine server at this point has the preprocessed data-points of wifi and 3G characteristics at different locations. Note that these are stored for all users, and we collaboratively use all users' data-points in the learning algorithms described in this section.

Intuitively, we want to look at the user's current movement patterns, correlate it with movement patterns in the past, and find the expected time-to-wifi. This is a prediction problem in continuous variables, and is a well-studied problem in machine learning and artificial intelligence. For this, we want to form a **prediction model**. This models the time-to-wifi as a **response** to (i.e. dependent on) some **covariates**. In our case, we model time-to-wifi as a response to covariates like location, bearing, speed etc.

More concretely, we are collecting data points, which correspond to a certain location, bearing, speed, and timestamp. We also measure the time-to-wifi corresponding to each data point by looking at wifi signals for 10 minutes after the initial data collection. Our goal is to form a prediction model which, given a new data point (location, bearing etc), can predict the time-to-wifi corresponding to it. To form this prediction model, we used various machine learning techniques, as discussed in the following section.

14

4.5   Regression

Since we are dealing with continuous variables, regression immediately springs to mind to form the prediction model.

Linear regression: Linear regression learns a linear model relating the response in terms of the covariates. It then uses this linear model to predict the response, given a new set of covariates. However, in our problem, it is expected that the wifi availability pattern at one location could be completely different from another location. Similarly, the time-to-wifi could be vastly different even in the same location for two different bearings. Therefore, it is not possible to have a simple linear relation between the response (time-to-wifi) and the various covariates like location and bearing. For linearity, we divided location and bearing, which were global variables, into sets of local variables instead. We did this by dividing the map into a grid of blocks based on the GPS coordinates, and further subdividing each block into **sectors** based on bearing. For example, one sector could include all the data points that have their GPS coordinates in the square bounded by $(40.345, -74.660) - (40.350, -74.655)$ and bearing in the range $45°$ - $90°$. This sector is shown in Figure 4.


Figure 4: Green points are one sector. They are the points in the bearing range $[0°, 45°]$ in the GPS coordinate square shown (bounded by $(40.345, -74.660) - (40.350, -74.655)$).

We then performed cross-validated error analysis to see what granularity of dividing latitude/longitude and bearing into sectors yields best results. Cross-validation basically forms the prediction model based on *part* of the dataset, and then uses this prediction model to predict time-to-wifi for the remaining part of the dataset. We can compare the predictions with the known values to find the error in prediction. By repeating this cross-validation for different sector sizes, we found that latitude/longitude divisions of $0.005°$ and bearing divisions of $45°$ work best.

Then, we modeled each data point as follows: each data point has some "global" covariates like speed, timestamp associated with it. As described above, the "global" variables latitude, longitude and bearing were split up into sectors. Therefore, each data point also has three covariates per sector of the map. If the data point falls in a sector, then the values for its 3

covariates for that sector would be the latitude, longitude and bearing difference of the data point from the minimum bounds of that sector. If the data point does not belong to a sector, its 3 covariates corresponding to that sector have value 0.

Then, using all the data-points from the past, we can come up with a linear model to predict time-to-wifi in terms of the many covariates.

Regularized regression: The same methodology as above was followed, except the model parameters/covariates were weighted via ridge regression. This allows us to penalize certain regression coefficients in a manner that reduces the variance in the prediction.
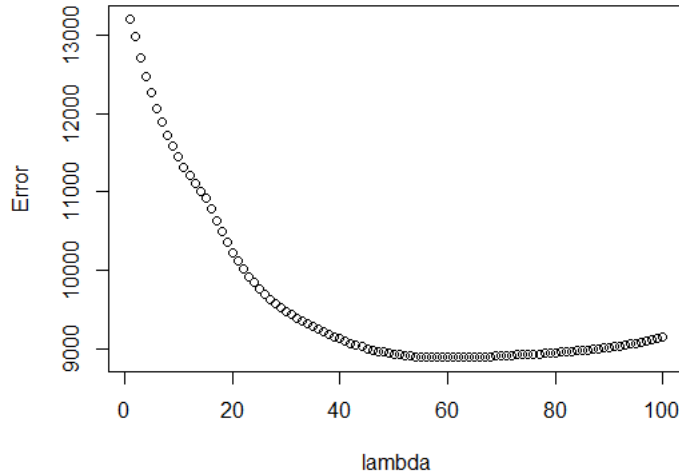

Figure 5: Cross-validated error vs. λ for ridge regression

To figure out the value of the ridge parameter λ, we found the cross-validated error for various values of λ (similar to method in linear regression). Figure 5 shows the plot of cross-validated error vs λ as obtained from R's `glmnet` package. Following the result of the graph, we used $\lambda = 60$ in our final analysis, since it shows minimum error.

4.6 k-means clustering

We use k-means clustering to form our prediction model. The general overview of the method is this: we use k-means to split the training dataset into clusters, based on only the covariates (not the response). Then, for predicting time-to-wifi, we use the same covariates to assign the new datapoint to a cluster, and set its time-to-wifi to be the mean time-to-wifi of all training data-points in the cluster.

In our case, we expect that two data points originating close to each other and showing movement in the same direction would obtain wifi after approximately the same time. Hence, the covariates we used for k-means clustering were latitude, longitude and bearing (note that speed was not considered in this method). We also scaled down bearing by a factor of 18000 in order to make the variation in bearing approximately equal to the variation in latitude and

16

longitude, since otherwise the clusters would form based simply on bearing (Note that across Princeton campus, latitude and longitude only changes by 0.02, compared to bearing which varies from 0° to 360°).

The time-to-wifi for each cluster was set to be the mean time-to-wifi of the data-points in that cluster. For predicting the time-to-wifi for a new data point, we find the cluster that it is closest to, in terms of latitude, longitude and bearing, and set its time-to-wifi to be the mean value for that cluster. Similar to the cross-validation in the previous sections, Figure 6 shows the plot of cross-validated error as a function of the number of clusters.


Figure 6: Mean cross-validated absolute error vs. number of clusters

4.7 Random forests

Once we realized that one of the major problems that we faced in terms of machine learning was sparseness of data (as described in Section 4.9), we looked for alternatives. One algorithm that we considered was random forests.

Random forests builds up a forest of *decision trees*. Decision trees are predictive models in themselves. They are a way of mapping the observations we see - location, speed, bearing - to the value that we want to learn - the time-to-wifi. Decision trees are best understood as following the path of making a decision based on the covariates. A part of such a decision tree is shown in the figure FIGURE

TODO: figure about decision tree

Therefore, a single decision tree can, given some value of the covariates, predict what the value of the time-to-wifi will be. However, the prediction depends entirely on the structure of the decision tree. Depending on the order in which we consider covariates going down the tree, or on whether we ignore certain covariates and emphasize others, we will get a different tree, and hence different predictions. This is where random forests comes in. It builds up a bunch of different decision trees using different splits of the covariates. Then, given a new data-point, it

17

uses all these decision trees to make various predictions of time-to-wifi, and returns the average value found.

We used random forests since it is quite effective with sparse data sets and it balances the error in datasets where the data are quite unbalanced, as is true in ours (again, further described in Section 4.9). This algorithm was quite easy to implement and analyze. We didn't need to restrict the covariates as we did in linear regression because random forests itself can find the best split of covariates to use to base its decisions on. It also gives a measure of how important each covariate is, which is useful in analysis (notably, it showed that the speed measurement was not very important).

4.8   Boosting

Boosting is a predictive method that uses several weak predictors in conjunction to form a relatively strong predictor. Consider the situation where we have a few predictors of time-to-wifi that don't perform well, i.e. they have quite a lot of errors. Boosting can use these together and improve the overall performance. It is a fascinating algorithm, since it is guaranteed to predict the real value in the limit, as we increase the number of predictors.

Boosting is something we tried more out of curiosity than anything. Since we had tried a few regression algorithms which had had a large error, we wondered if using boosting would help bring down the error by using multiple weak regression algorithms in conjunction. However, we found that boosting did not result in much better predictions, and in fact, in some cases it actually did poorer.

4.9   Testing and Analysis

To test the value of using each of these algorithms, we used 5-fold cross validation. For this, first we split the training dataset into 5 parts. We train the prediction model on 80% of the data. Using this model, we predict the time-to-wifi for the remaining 20% of the data. Since the actual time-to-wifi for this 20% is known, we can obtain the mean absolute error in prediction of time to wifi. The blue bars in Figure 8 show the results for each of the 5 algorithms. While random forest regression seems to perform significantly better overall than the rest, the errors are fairly high (35-70 seconds). We hypothesized that these errors could be due to the outliers in our data. These outliers could occur for two reasons:

1. *pathological movement patterns*: Firstly, users' walking patterns can be non-uniform. For example, most people heading towards the CS building from the Shapiro Walk will obtain wifi in less than a minute. However, say a user who is taking this path suddenly remembers they have to go somewhere else, and end up turning away. The application will end up recording that the time-to-wifi along this path is much longer than it actually is. Consider what would happen if we were making our prediction model using k-means. These "peculiar" data points, if clustered along with the more regular ones, can have predictions vastly different from the observed responses and hence skew our measure of average error.

   This source of error is more or less an unavoidable property of the problem we are trying to solve. The best we can do is to go ahead determine how long to delay the app under the given prediction model.. If, during this delaying period, the application does not get wifi connectivity (perhaps due to misprediction), we can send the apps' network traffic over 3G instead.

18

2. *sparse data*: Secondly, as can be verified from Figure 3, certain areas have points that are more spatially separated. This sparseness of data means that we may not have enough information about the various possible movement patterns that a user could be walking along. This makes prediction hard, and hence can lead to high errors under cross-validation.
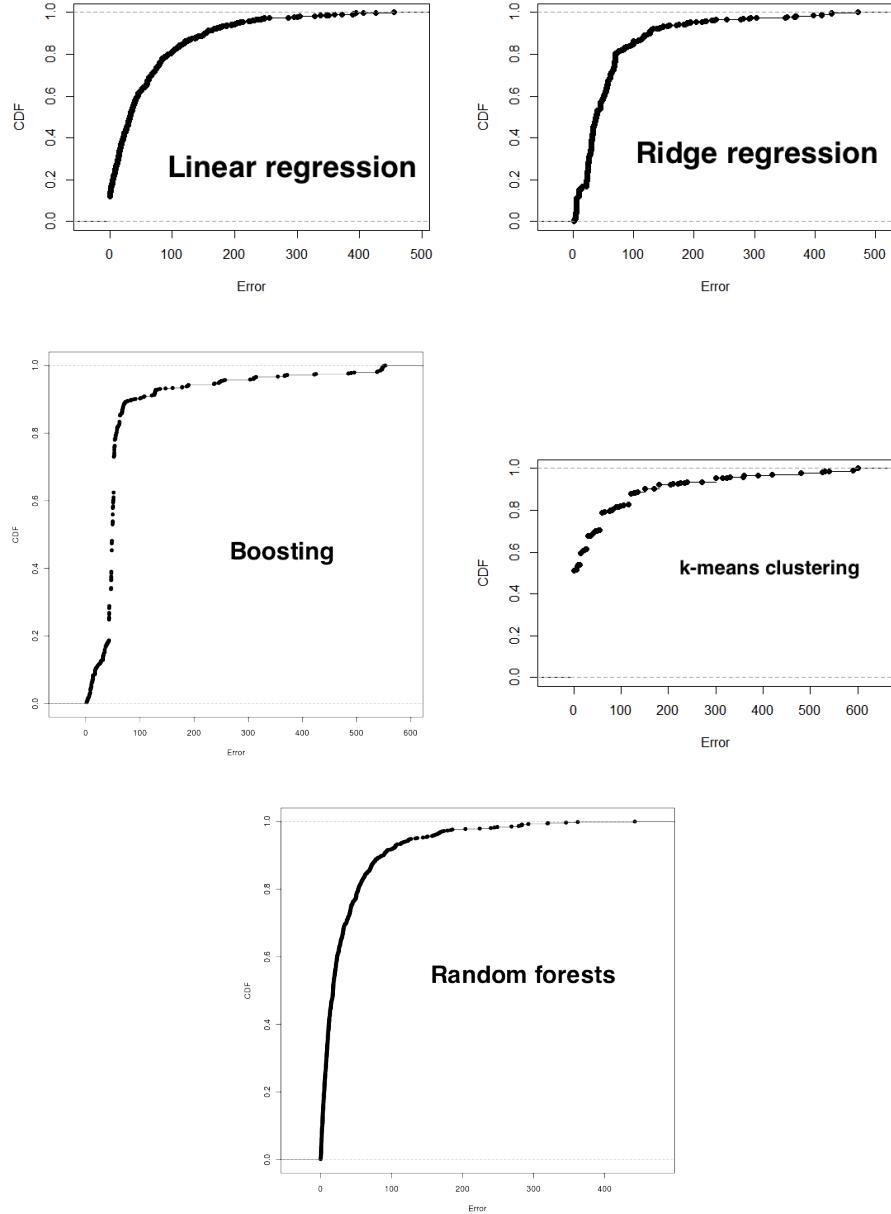

Figure 7: CDF of mean prediction errors for different algorithms

To find out if sparseness of data is an issue, I looked at the cumulative distribution (CDF) of the prediction error made by the different algorithms. This is reflected in Figure 7. These CDFs show that the errors are more heavily weighted towards small errors, i.e. a large proportion of prediction errors are small. This is best seen in the CDF for random forests, where we can see that $\approx 80\%$ of the errors are on the order of 15 seconds.

In order to further verify that average error was being skewed by a few large errors, we re-ran the algorithms, focusing only on the part of the campus where we managed to obtain a high density of data points (area around Friend Center, CS Building, ORFE, E-Quad, Shapiro Walk). Figure 8 shows the new results. Just as we expected, the average error decreases significantly to almost half of before, with all algorithms performing similarly well. This result indicates that we do not have enough data that is distributed well over the entire campus. However, we believe that time-to-wifi estimator has significant value when it makes errors of ≤ 1 minute, which happens most of the time. Further, the predictions only seem to get better with increased density, and hence it is likely that we will get better results as more people use the system.


Figure 8: Mean cross validated absolute error for each algorithm (with and without outliers)

4.10 Prediction on the phone

time-to-wifi-predict

Now, we have used the data-points on the server to come up with a predictive model. We need to get the predictive model to the phone, and set up a way of using the model to predict time-to-wifi. For this, we firstly need a convenient way to represent a predictive model. This is quite easy to do, and in fact, most machine learning packages provide a way to output the predictive model in an easy-to-use format. For example, for k-means clustering, we only need to keep track of the centers of the clusters, and their corresponding time-to-wifi values.

We can then co-opt the same data collection app described above (Section 4.1) to receive the predictive model from the server, using a simple HTTP request to fetch a `.txt` file. For

space-efficiency, we only send that part of the predictive model which corresponds to the user's current location and nearby areas (since the predictive model for the entire world can be quite large). This app exposes this predictive model through a method which can be called to predict the time-to-wifi, given the user's current location, bearing etc.

In summary, the app collects the data for learning the predictive model, and sends this data to a server. Machine learning algorithms run on the server to form a predictive model, which is sent back to the same app. This app then acts as the time-to-wifi predictor through an interface with the rest of the applications.

4.11 Future Improvements

In the future, the following improvements could be made to the time-to-wifi predictor:

- In making the time-to-wifi predictions, we use all users' data-points of wifi and 3G locations, and we consider each user's data to be equal. Maybe in the future we could weight the current user's data a little more in making predictions, since it more likely to reflect the user's movement patterns.

- The current solution requires machine learning to be done offline on a server - this was because it would be easier to program and also faster. However, a more realistic solution would probably use some greedy online heuristics on the phone to learn time-to-wifi on the fly as network requests are being made (instead of needing a pre-computed prediction model).

- For measuring time-to-wifi, we are currently using wifi signal strength as an indication of how usable the wifi signal at a given point is. However, various papers have shown that signal strength is a poor predictor of wifi quality. We currently use it in the system out of ease of programming. Ideally, however, we should test wifi quality by doing some sort of bandwidth test. This could involve sending small packets of data to ping a known server and waiting for a response. This empirical method would help us more accurately gauge the quality of wifi.

5 Application delay tolerance estimator

- Data collection - app screenshot

- Current solution

- Results

- Future improvements: Not just poll every so often, but update traffic stats every network request (disadv: requires kernel access?)

The other input into the policy is the application delay tolerance estimator. Just to recap, this estimator will learn from the apps' past network usage statistics - how much data does each app send, and what is the usual pattern of this data generated. We expect, for example, that a video streaming app would be show constant, high-bandwidth pattern, as contrasted to a web browser where we would expect more bursty, low size traffic as the user clicks around on links.

Then, when we see a certain app making network requests we can look at its past network patterns to gauge its delay tolerance. This will be an input into the policy.
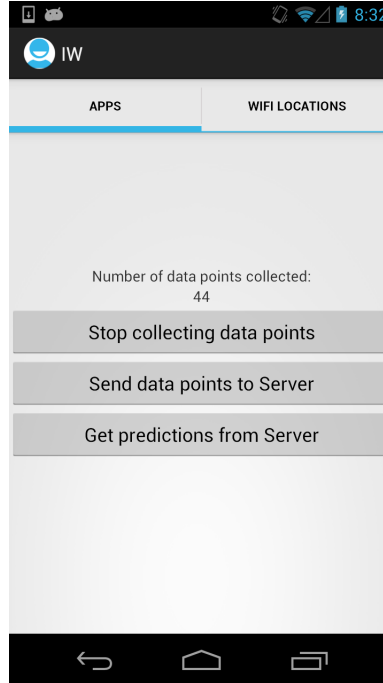
5.1 Data collection

Figure 9: App for collecting network usage statistics of other apps

For the delay tolerance estimator, we first need to collect data about the network usage of all apps. For this, I built another app that runs on the phone in the background and records apps making network requests, storing the network usage statistics of each app at different times.

This app uses the `TrafficStats` class which exposes the following four different network data statistics for each app:

1. number of bytes sent over TCP

2. number of bytes sent over UDP

3. number of bytes received over TCP

4. number of bytes received over UDP

The app periodically polls the `TrafficStats` class and records the above values for all the apps at that time. NB: if an app's network stats have not changed at all since the previous recording, meaning that it has had no network activity in the intervening time, then we don't record its value again, to avoid redundancy.

We can see that this quantizes the network traffic pattern graph. If we make the first record from `TrafficStats` at time t_1 and then the second at t_2, then we are attributing all the network traffic in the interval $(t_1, t_2]$ to time t_2. However, if our `TrafficStats` sampling is done frequently

enough, then we will approach a quantization level where it will be imperceptible to the user. We chose a sampling rate of 100 ms, which is fine because we will be calculating delay tolerance in the order of seconds and so this is an order of magnitude better resolution.

5.2 Learning delay tolerance

In trying to learn the delay tolerance of apps, the question arose as to whether to do the machine learning on the phone rather than the server. Initially, we attempted to follow a method similar to that used to find the time-to-wifi - doing offline learning on the server to form a compact prediction model on the phone. While this method seemed more rigorous, it was quite hard to implement algorithms to learn delay tolerance. However, a key insight showed us that the learning of delay tolerance could be a matter of simple statistical analysis, and hence could be done on the phone. We can sample all the network statistics collected by the app described above, and analyze these in an *online* manner. This means that the amount of data stored is always a constant amount, and we need not make expensive requests to send data-points to the server, and receive prediction models from the server.

So how do we actually do all this on the phone? We now know the network statistics of all apps running on the phone. Using this information, we want to determine the apps' delay tolerance. When a user is directly interacting with an app, any network traffic generated is not delay tolerant. This is obvious - we cannot delay traffic that the user is explicitly waiting for without adversely affecting their interaction and experience with the app. Therefore, we discard statistics about these situations. From the remaining stats, we can find out the GET requests made by the app in the *background*. For applications like email clients and messengers, these background requests tend to be requests for background syncing. As we have already discussed at length, such network requests are usually tolerant to delay as long as the user isn't directly interacting with the app making the request.
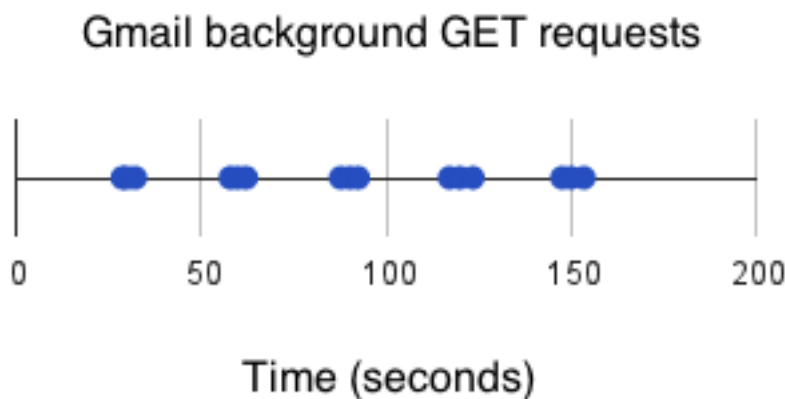

Figure 10: Background GET requests sent by Gmail app

So, now that we know that what sort of traffic we can actually delay, we want to be able to figure out how long we can delay the network traffic for, i.e. what is the delay tolerance of the traffic. Figure 10 shows background GET requests for the Gmail client, whose background traffic I logged. It shows that the app sends out syncing requests every so often as a *cluster* of 3-4 separate GET requests. Note that these multiple near-simultaneous requests are not retransmissions.

Rather, they all serve different purposes: for e.g. on the email app, these are multiple concurrent requests to refresh the various different email folders. Also note that the fact that we are logging network statistics only every 100 milliseconds or so does not affect these `GET` requests. This is because the `GET` requests are separated by seconds rather than milliseconds.

The time between these successive requests is an indication of how much the app is willing to wait before pinging for updates. This is in line with our definition of delay tolerance: if the app is willing to wait one minute before sending a background sync request, then that must mean that it's background network requests are tolerant to delays on the order of one minute. Note that this pattern seems to be true in the background usage statistics of other apps as well - I looked at K9, an open-source email app, and Facebook as well. ¡DRAW GRAPH FOR THESE APPS¿

In the context of the `GET` request graph shown above, we can come up with a method for actually gauging delay tolerance. To learn delay tolerance of an app, we should look at the time periods between the clusters of requests. This means that we should consider the clusters of concurrent requests as one single request. Naively, this is quite simple to do. We can simply store the timestamp of the last `GET` request sent by the app. Then, when a new background `GET` request is sent, we can subtract the stored timestamp value from the current time to get an estimation of delay tolerance.

However, there are improvements which can be made to this method. First of all, we can see that we are recalculating delay tolerance every time, and we don't end up making full use of the history of the app's network usage. If we simply follow the method above of subtracting the past timestamp from the current, we have a model that is only looking at the past two time steps and discarding the rest of the information about the past. Rather than doing so, we can calculate delay tolerance as a time-weighted average of the time differences we are recording. Mathematically, we can define this recursively as follows:

$$D_i = \lambda \Delta t_i + (1 - \lambda)D_{i-1} \tag{1}$$

where λ is the weight for taking the average (the higher the value of λ, the more we weigh recent observations as compared to past observations). To unpack this equation a bit more, we are calculating the i^{th} estimate of delay tolerance using the weighted average of 1) the $(i-1)^{\text{th}}$ estimate of delay tolerance, and 2) the time between the two most recent `GET` requests.

So far, we have assumed that `GET` requests are infrequent, periodic, single pings. However, in reality as we saw above, apps tend to send out a cluster of separate `GET` requests close together. We need to be able to account for this. Intuitively, this will require us to ignore or discount `GET` requests if another `GET` request was made by the same app very recently. If we choose to *discount* them, then we can simply have a threshold on the amount one `GET` request can change the delay tolerance estimate. For example, we can say that a single request can only change the delay tolerance by at most $\pm 10\%$, and clamp any bigger changes to this maximum value. On the other hand, we can also *ignore* `GET` requests that come too close together. Therefore, if we see many requests coming together, we can ignore all but one of them, i.e. we can cluster them together. To do this, we can see when two `GET` requests are closer together than a minimum fraction of the current delay tolerance estimate, i.e. if

$$\frac{\Delta t_i}{D_{i-1}} < \epsilon \tag{2}$$

we can ignore Δt_i in the delay tolerance calculation.

In summary therefore, we have uncovered three (related) methods to gauge delay tolerance:

1. *Basic method:* Simply take the weighted average of current delay tolerance estimate and the time between last two `GET` requests (see Equation 1)

2. *Clamping* the change in delay tolerance estimates to some maximum factor, say 10%.

3. *Clustering* requests that are close together by ignoring all but one of them.

Note that these are online algorithms, i.e. they examine the stream of data as it comes in, but only store a constant amount in order to make their estimates.
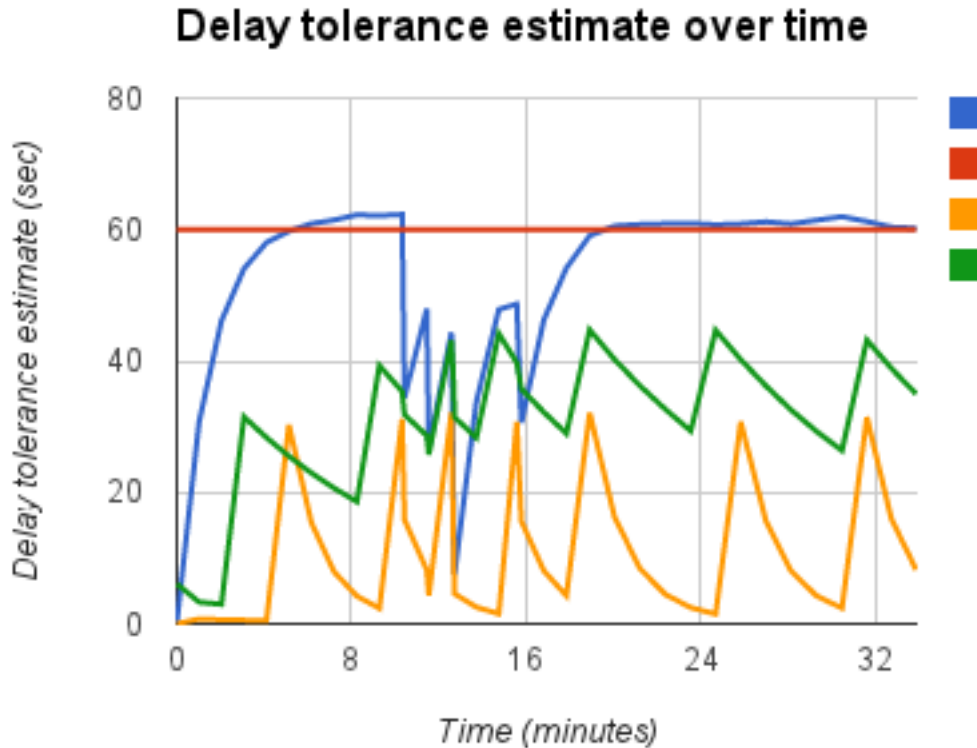
5.3   Results and Analysis

Figure 11: Estimation of delay tolerance

We now move to discussing the relative merits and demerits of the three methods of learning delay tolerance. Above is a graph which shows the delay tolerance estimates made over time by the three methods, compared with the real delay tolerance of 60 seconds. These measurements were made for the open-source K9 email client.

1. The basic method (orange line in Figure 11) is a naive method. It assumes that background sync occurs through a single network request. However, in reality, quite a few packets are sent together. Since these have very short intervals between them, it brings down the delay tolerance estimate greatly. This is reflected in the valleys in the above graph, where the delay tolerance estimate goes down quickly due to many `GET` requests. The steep climb thereafter is due to the fact that the app waits and doesn't send any `GET` requests after sending out sync requests for the first time.

   However, it is interesting to note that, as a result of this method's naivete, it ends up considering all `GET` requests, and so it is guaranteed to always be bounded from above by the true frequency of background syncing (the red line). Therefore, we are always underestimating delay tolerance and erring on the side of caution. On the other hand, we are not able to delay packets as much since the delay tolerance is lowered.

2. The method of clamping (green line in the graph) helps reduce the effects of multiple requests coming in together. Since each of these requests can only change the delay tolerance estimate by at most 10%, we don't see as much of a drop as we saw with the basic method.

   However, even 4 requests coming in together would result in a drop down to 65% of the original value, and as a result, we can see that this method also underestimates the delay tolerance value, albeit not as much as the basic method.

3. Instead, if we just cluster the concurrent requests together, we get the red line in the graph above. This method does much better at predicting delay tolerance. Since it considers only the time between subsequent syncing requests, it provides a good approximation of the real frequency of background syncs. As we can see in the graph, the estimate of delay tolerance stabilizes around the value of 60 seconds.

   Note that this method also has its pitfalls. It depends on a good selection of ϵ as defined in Equation 2. If the value of ϵ is too low, we will end up considering `GET` requests that we should have ignored. In fact, if we look at Figure 11, this is precisely what we see from $x \approx 10$ to $x \approx 20$. Here, the `GET` requests came close together, but not close enough to be clustered together according to Equation 2. Therefore, we end up getting a low estimate of delay tolerance before it settles back to the real value.

5.4   Future work

The following improvements could be made in the future to the delay tolerance estimator:

- As we have described above, the data collection app quantizes time. It only polls for changes in network stats every 100 milliseconds or so, and therefore, ends up having a coarse-grained idea of when apps actually send and receive data. Instead, if we could actually look at the network requests as they happen, we would have a better estimate of delay tolerance. Our initial reservation against doing this was over the fact that it would require working in the kernel in the network stack. However, as described in Section 6, I had to fiddle with the network stack anyway for another purpose. Therefore, it should now be easy to implement the data collection app for the delay tolerance estimator in the kernel and get more precise measurements.

6 Mechanism

At this point, we have built the 2 machine learning components which provide the time-to-wifi and delay tolerance estimates. Using these, we can construct a policy which will inform us about how much to delay network traffic. However, before we delve into this, let us consider a mechanism for delaying traffic, i.e. some actual code that runs on the phone to delay network packets for some requisite amount of time. As explained in Section 2.5, we want to implement the mechanism in the network stack in the kernel, so as to avoid requiring application developers from changing their apps in any way.

Therefore, the general approach is to insert ourselves in the network stack, so that we can examine all network requests that are being made. If we find that wifi is not available when trying to make a network request, we can ask the policy whether that request can be delayed, and if so, how long it can be delayed for. Since we are in the network stack at this point, we can delay the network request in the hopes of getting wifi soon, and if successful, we can send the request out over wifi. However, if the delaying timer expires without us getting in range of wifi, we will just send the request over 3G so as to not affect the user experience.

6.1 Serval

Serval is a project of a networking group at Princeton University led by Michael Freedman. The project aims to modify the network stack to allow for service-centric networking, i.e. networking set up around network "flows" between the user and the service they are using. These flows account for dynamism (e.g. due to user mobility) and multiplicity (e.g. due to multiple interfaces at the user end) in the networks underlying these services. This was not previously possible in the traditional host-centric TCP/IP stack. They introduce a Service Access Layer (SAL) on top of this traditional network stack, which allows applications to interact over service names (instead of hostnames).

As a part of their modifications to the network stack, they introduced mechanisms to allow for migration of network flows from one interface to another. We can see how this would be useful for service-centric networking - it allows a flow to continue uninterrupted even if the user is moving around and switching from wifi to 3G, or between various wifi access points. We can co-opt this mechanism in Serval to allow us to switch between wifi and 3G, for example if we delay a network request that was going to be sent over 3G and end up getting wifi before the delay timer expired.

Ozlem Biglir Yetim, as part of her own research described in Section 3, introduced another mechanism into Serval which allowed for the delaying of network flows. This is again useful for our work, since we need to delay network requests according to the policy if wifi is not available.

Because of these delaying and migration mechanisms, and also because working with the SAL codebase would allow me to immediately enter the network stack, I decided to use Serval to implement the mechanism. It was relatively straightforward to modify SAL to allow us to look at every network request. If wifi wasn't available, we could use the policy to figure out how long to delay the request, and delay it accordingly using functions already existing in the codebase.
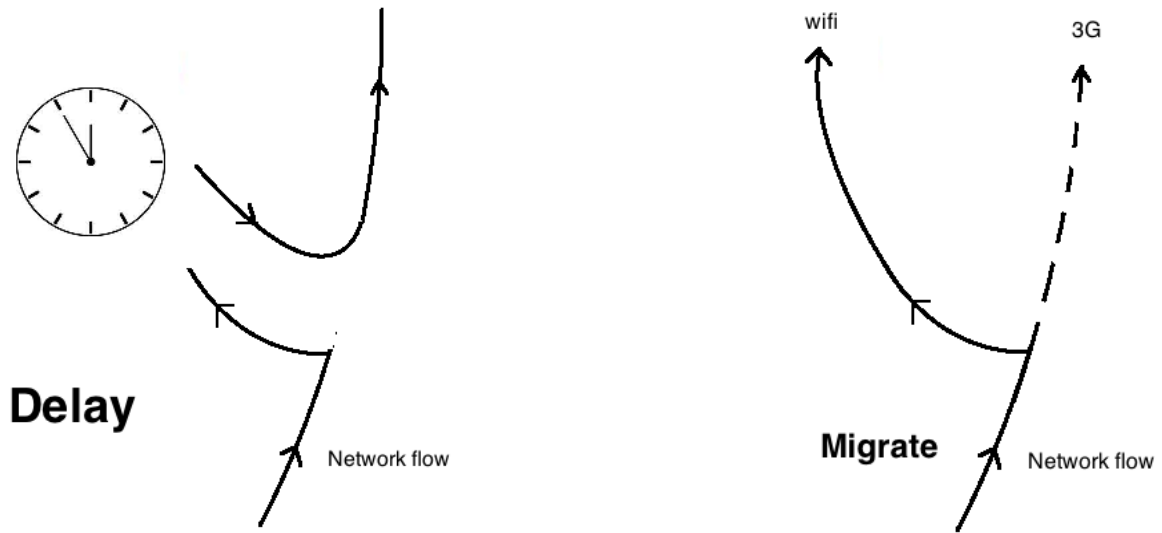
Figure 12: Serval can be used for delaying network flows and migrating them across interfaces.

7 Overall policy

- "The goal is to compare the performance of different policies in different settings (for example, settings where wifi and 3G are geographically well mapped out, as compared to settings where they are not)."

Plan:

- Just walk around with the phone with k-9 running, with 3G and wifi, and see how many delays actually happen, and how many are successful.

- Log this in a way such that it is policy independent. For example, you can record when the delay prediction was made, how much it was delayed, and how much time it actually took to wifi (note that if we don't get wifi by the time the delay tolerance runs out, we switch to 3G always) - this method of recording data will allow us to analyze different policies using the same dataset, e.g. a policy which delays if time-to-wifi ¡ delay tolerance, or a policy which delays for fixed 10 sec can both be tested using this dataset.

- Try to come up with different policies

- As you're finishing up, try to look for papers for each of the points - e.g. wifi power level measurement using signal strength, papers on wifi vs 3G battery consumption. Basically look for previous work on anything that I've done, like using learning to predict movement patterns (specifically Markov models).

8 Experiments

9 Results

10 Conclusions

References

[1] Ananthanarayanan, Ganesh, and Ion Stoica, "Blue-Fi: Enhancing Wi-Fi Performance using Bluetooth Signals." *MobiSys '09*, Krakow, Poland, June 22-25, 2009: p. 249-261.

[2] Android developer guide, <http://developer.android.com/reference/packages.html>.

[3] ATT 3G coverage in USA, <http://www.wireless.att.com/coverageviewer/?type=data>.

[4] ATT data plans, <http://www.att.com/shop/wireless/data-plans.html>.

[5] Balasubramanian, Niranjan, Aruna Balasubramanian, and Arun Venkataramani, "Energy Consumption in Mobile Phones: A Measurement Study and Implications for Network Applications." *IMC '09*, Chicago, USA, November 4-6, 2009: p. 280-293

[6] Ericsson, "Ericsson Mobility Report." Stockhold, Sweden, February, 2013.

[7] GSM Association, "Universal Access Report." London, UK, 2006.

[8] Han, Dongsu, and Michael Kaminsky, "Getting the "Inside " Scoop on Neighborhood Networks." *IMC'08*, Vouliagmeni, Greece, October 20–22, 2008.

[9] Lee, Kyunghan, Joohyun Lee, Yung Yi, Injong Rhee, and Song Chong, "Mobile Data Offloading: How Much Can wifi Deliver?" *IEEE/ACM Transactions on Networking*, Philadelphia, USA, June 2012.

[10] Ra, Moo-Ryong, Jeongyeup Paek, Abhishek Sharma, Ramesh Govindan, Martin Krieger, and Michael Neely, "Energy-Delay Tradeoffs in Smartphone Applications." *MobiSys '10*, San Francisco, USA, June 15-18, 2010: p. 255-269.

[11] Rahmati, Ahmad, and Lin Zhong, "Context-for-Wireless: Context-Sensitive Energy-Efficient Wireless Data Transfer." *MobiSys '07*, San Juan, Puerto Rico, USA, June 11-14, 2007.

[12] Raman, Bhaskaran, and Kameswari Chebrolu, "Experiences in Using WiFi for Rural Internet in India." *IEEE Communications Magazine*, January 2007: p. 104-110.

[13] Verizon 3G coverage in USA, <http://opensignal.com/network-coverage-maps/verizon-coverage-map.php>.

[14] Verizon data plans, <http://www.verizonwireless.com/wcms/consumer/shop/share-everything.html>.

[15] Yetim, Ozlem Bilgir, and Margaret Martonosi, "Adaptive Usage of Cellular and wifi Bandwidth: An Optimal Scheduling Formulation." *CHANTS '12*, Istanbul, Turkey, August 22, 2012.